🟠 Registration Now Closed for Enterprise 2026

ENTERPRISE 2026 · EVALUATION OVERVIEW

THE LEADING INDEPENDENT SECURITY EVALUATION

ATT&CK Evaluations Enterprise 2026 tests how your solution detects and stops modern adversary behaviors — using real tradecraft, a unified scoring framework, and no vendor-controlled results. Registration is now closed.

- New: Total Evaluation Score (TES)
- EDR + XDR + MDR + AI SOC unified
- Results published December 2026
- Application deadline April 30, 2026

Participation Information | Methodology Overview | Citation Policy

WHY PARTICIPATE

Join the leading independent evaluation of enterprise security solutions

ATT&CK Evaluations are different from vendor-run tests, analyst reports, or ad-hoc pentests. They're the only independent, behavior-based evaluation that tests how solutions detect and stop real adversary tradecraft — with full telemetry published openly.

Credible Third-Party Validation

MITRE is the most trusted name in threat intelligence. An ATT&CK Evaluations result is recognized by CISOs and security teams worldwide as independent and unbiased evidence of product quality.

🔍

Real Gap Identification

Per-technique DC and PC results show your engineering team exactly where coverage gaps exist — against Critical and High-priority techniques, not just overall averages. The most actionable product roadmap data available.

🏆

Differentiated in the Market

Enterprise 2026 introduces TES — a unified quantitative score — for the first time. Early adopters of the new framework establish credibility with buyers before the methodology becomes the market standard.

💡 **Note on competitive context:** ATT&CK Evaluations do not rank vendors or declare winners. Results show what was achieved against specific adversary scenarios and how. Customers use the full breakdown — TES, modifier, per-technique results — to assess fit for their environment. Participation demonstrates commitment to transparency.

NEW FOR 2026

What's New in Enterprise 2026

Total Evaluation Score (TES) Framework

For the first time, ATT&CK Evaluations introduces a unified quantitative scoring system. TES combines detection quality (DQI) and protection quality (PQI) into a single comparable metric on a 0–2.0 scale. This moves beyond binary "detected / not detected" results to measure alert quality, analyst precision, platform speed, block timing, and false positive performance — all weighted by technique criticality.

Transparent Source-of-Action Attribution

Every TES score is published with a modifier showing who or what produced the quality:

MODIFIER	WHO OPERATES IT	OPERATIONAL MODEL
(P) Platform	Fully automated platform	Customer SOC team investigates from automated alerts
(P/AI) Platform + AI	AI-augmented automation	AI enriches alerts; analysts make final decisions
(S/H) Service + Humans	MDR/MSSP human analysts	Analysts validate, enrich, and escalate complete incidents
(S/AI) Service + AI SOC	Autonomous AI operations	AI performs analyst-equivalent investigation and recommendations
Mixed	Hybrid	Combination of platform automation, human analysts, and AI

Unified Platform and Managed Services Evaluation

Enterprise 2026 eliminates the historical separation between "Enterprise" and "Managed Services" rounds. EDR, XDR, MDR, MSSP, SIEM, and AI SOC vendors are all evaluated on the same adversary scenarios using the same quality metrics. Source-of-Action modifiers provide execution transparency without requiring separate tracks.

Vendor Tooling Declaration

All participating vendors must declare product category tags ensuring customers understand exactly what was tested. Tags: EDR, XDR, Cloud Security, Identity Protection, MDR Service, Network Detection, Managed Service / MSSP, SIEM Integration, Endpoint Protection. Multiple tags are permitted.

EVALUATION SCENARIOS

Threat Context & Scenario Design

Threat Overview

The 2026 ATT&CK® Evaluations reflect the current enterprise threat landscape, where financially motivated intrusions dominate overall activity and state-sponsored operations continue to pose significant risk — particularly in cloud and infrastructure-focused environments.

operations. Effective detection and response therefore depend on behavioral context, correlation across systems, and understanding attacker intent — not isolated indicators.

💰 **Financially Motivated**

Crimeware-as-a-service, double-extortion operations, credential abuse — high-tempo intrusions dominating enterprise incident response

🏛️ **State-Sponsored**

PRC cyber espionage operations targeting diversified infrastructure — persistent, deliberate, identity-based lateral movement across hybrid environments

🎭 **Shared Tradecraft**

Both threat types increasingly use living-off-the-land, valid credentials, and legitimate tools — making behavioral correlation the essential detection primitive

Scenario Representation

The two scenarios together reflect the prevailing threat balance — the majority of enterprise intrusions are cybercriminal in origin, while state-sponsored actors continue to drive innovation in infrastructure exploitation.

SCENARIO 1 · ENDPOINT FOCUSED

Financially Motivated Tradecraft

● Registration closed

Crimeware-as-a-Service Operation

This scenario reflects the layered, role-based economy of cybercriminal operations, where specialized actors collaborate across distinct phases of an intrusion without operating as a unified team. It begins with targeted social engineering to harvest credentials, which are then monetized and passed to downstream actors who leverage that access to compromise victim environments and execute on their objectives. Each role in the chain is optimized for results, reflecting how the commoditization of cybercrime has made intrusions faster, harder to attribute, and more accessible to a broader range of adversaries.

Scenario Focus: An endpoint-focused intrusion tracing the crimeware-as-a-service lifecycle, from social engineering-driven credential theft, to environment compromise, lateral movement, data exfiltration, and destructive impact, with living-off-the-land techniques employed throughout to evade detection.

Ideal for: EDR platforms, SIEM integrations, endpoint protection vendors, managed services focused on host-based security

● Registration closed

PRC Cyber Espionage Group

A technically advanced PRC cyber-espionage actor that poses a persistent threat to organizations worldwide. The group has operationalized zero-day and n-day vulnerabilities, and exploited internet-facing edge devices and remote access infrastructure to gain footholds. Environmental adaptability, custom payload development, living-off-the-land behaviors, and legitimate cloud services are leveraged to blend into enterprise activity. The group's tradecraft broadly represents how PRC state-sponsored actors operate, sharing tools, techniques, and infrastructure across a well-resourced offensive cyber ecosystem. The convergence complicates attribution and makes the threat relevant to organizations across sectors and geographies.

Scenario Focus: An enterprise-wide, evasive intrusion featuring exploitation of diversified infrastructure, deployment of custom and open-source tooling, and abuse of legitimate cloud services.

Ideal for: XDR platforms, identity protection solutions, MDR/MSSP providers, enterprise-wide visibility offerings

Evaluation Intent

ATT&CK Evaluations are designed to assess how security platforms and services detect, correlate, and respond to modern, trust-abusing behaviors. The evaluation emphasizes three core capabilities:

BEHAVIOR-BASED DETECTION

Identifying adversary activity from behavioral patterns rather than signatures or isolated indicators — what the adversary *did*, not just which tool they used

MULTI-EVENT CORRELATION

Connecting related events across time, systems, and domains into coherent attack narratives — reducing analyst burden while improving detection fidelity

SIGNAL VS. NOISE DISCRIMINATION

Distinguishing malicious activity from legitimate administrative or user behavior — especially critical where adversaries deliberately blend into normal operations

Technique Scope

The Technique Scope defines which ATT&CK techniques and sub-techniques are in scope for evaluation and how parent-child relationships are handled during scoring. This structure ensures precise, actionable detections and avoids ambiguity when evaluating security controls against adversary behaviors.

Parent techniques	Encompass additional adversary behaviors, implementations, and variations beyond what is captured by any specific sub-technique. When a parent is in scope, the evaluation assesses the broader behavior category.
Sub-techniques	Provide a more specific, detailed breakdown of particular adversary methods. When a sub-technique is in scope, scoring is evaluated against that specific implementation.
Scope boundary rule	If an adversary behavior does not fit neatly into a specific sub-technique but aligns with the general concept of the parent, it remains in scope under the parent technique. This prevents adversaries from evading evaluation scope through minor implementation variations.

🎉 **April 7, 2026 — Technique Reference Release.** This section lists techniques and sub-techniques in scope for reference. Final scenario scope is a subset of these techniques. ATT&CK Enterprise v18.1

SCENARIO 1 | SCENARIO 2

Search by ID or name...

67 techniques · 13 tactics

VIEW IN ATT&CK NAVIGATOR ↗ | DOWNLOAD JSON | OPEN IN ATT&CK ↗

ALL TACTICS (67)

- Collection 5
- Command and Control 9
- Credential Access 3
- Defense Evasion 11
- Discovery 17
- Execution 3
- Exfiltration 6
- Impact 1
- Initial Access 3
- Lateral Movement 1
- Persistence 5
- Privilege Escalation 3
- Reconnaissance 0

COLLECTION — 5 techniques

T1005	Data from Local System	↗
T1039	Data from Network Shared Drive	↗
T1074.001	Data Staged: Local Data Staging	↗
T1074.002	Data Staged: Remote Data Staging	↗
T1114.001	Email Collection: Local Email Collection	↗
T1114.002	Email Collection: Remote Email Collection	↗
T1119	Automated Collection	↗
T1213.003	Data from Information Repositories: Code Repositories	↗
T1213.006	Data from Information Repositories: Databases	↗
T1530	Data from Cloud Storage	↗
▶ 3 T1560	Archive Collected Data	↗

COMMAND AND CONTROL — 9 techniques

Vendors select Scenario 1, Scenario 2, or both based on product capabilities and target markets. Scoring is independent per scenario.

EVALUATION STRUCTURE

What Gets Evaluated

Detections Testing

Protections are disabled so the full adversary attack chain executes. Your solution demonstrates complete visibility. Evaluated on:

DIMENSION	WHAT'S MEASURED	SCORING COMPONENT
Detection Coverage	Alert quality per ATT&CK technique — does your alert provide technique-specific details, ATT&CK mapping, and all 7 context elements (WHO, WHAT, WHEN, WHERE, HOW, WHY, Recommended Action)?	DC tier → Weighted_DC → DQI
Detection Precision	Analyst efficiency — finding real threats without false positives, consolidating related activity into actionable cases rather than fragmented alerts	DP → DQI
Detection Speed	Elapsed time from technique execution to first automated platform alert. Human analyst escalation time is excluded.	DS → DQI

Protections Testing

Protections are enabled. Tests are stage-gated: once you block, the test ends. Evaluated on:

DIMENSION	WHAT'S MEASURED	SCORING COMPONENT
Protection Coverage	Block timing in the attack chain (primary) + context quality of the block alert (secondary). Blocking before the first critical stage scores PC-3; blocking at the terminal stage after all critical objectives scores PC-1.	PC tier → Weighted_PC → PQI
Protection Precision	False positive validation: can you block threats without disrupting legitimate business activity? Each incorrectly blocked benign activity	PP → PQI

reduces your PT score.

False Positive Validation

Both scenarios include dedicated false positive testing with legitimate business activities — IT administration scripts, approved software downloads, remote desktop connections, cloud storage access, network file transfers, and data processing tools. Solutions that over-block benign activity receive lower precision scores that feed directly into PQI.

INFRASTRUCTURE

Test Environment

Enterprise 2026 runs across two separate ranges — one Detection range and one Protection range that covers both scenarios. Both ranges are deployed together as a part of the Core Baseline Infrastructure (CBI), ensuring fair, repeatable testing across all participants.

Detections Range SCENARIOS 1 & 2

Both detection scenarios run on its own dedicated range. Protections are disabled during detection testing to ensure the full attack chain executes and all adversary behaviors are observable.

Scenario 1 — Endpoint Focus (Detections)

COMPONENT	SPECIFICATION
Domain Controllers & File Servers	Windows Server 2022
Web Servers	Windows Server 2022
Workstations	Windows 11 Enterprise
Email Infrastructure	Exchange Server 2019
Network Firewall	pfSense or Open Source
Network Segmentation	Realistic segmentation — DMZ, internal workstation zones, server zones

Scenario 2 — Enterprise Focus (Detections)

Includes all Scenario 1 components plus the following extended enterprise attack surface:

components	· Exchange Server · pfSense firewall · network segmentation
Cloud Infrastructure	AWS (Amazon Web Services)
Virtual Machines	AWS instances and/or additional virtualization platforms
SSO / Identity Provider	Authentik (self-hosted IdP) + AWS IAM Identity Center — federated across both AWS organizations
Edge Devices & Appliances	Network edge devices and appliances
Linux Systems	Ubuntu 24.04 (Noble Numbat)

Protections Range

SCENARIOS 1 & 2

Both the Endpoint and Enterprise scenarios are executed on a single shared Protections range. This range contains the full combined component set spanning both scenarios.

COMPONENT	SPECIFICATION
Domain Controllers & File Servers	Windows Server 2022
Web Servers	Windows Server 2022
Workstations	Windows 11 Enterprise
Email Infrastructure	Exchange Server 2019
Network Firewall	pfSense / Open Source
Network Segmentation	Realistic segmentation — DMZ, internal workstation zones, server zones
Cloud Infrastructure	AWS (Amazon Web Services)
Virtual Machines	AWS instances and/or additional virtualization platforms
SSO / Identity Provider	Authentik (self-hosted IdP) + AWS IAM Identity Center — federated across both AWS organizations
Edge Devices & Appliances	Network edge devices and appliances
Linux Systems	Ubuntu 24.04 (Noble Numbat)

Participant Access & Integration

CLOUD COMPONENTS

AWS Victim Account — IdP-Federated Access

Participants receive account-level access to the victim AWS account(s) representing the cloud environment.

- ✓ Access is granted through **Authentik** (self-hosted IdP) federated with **AWS IAM Identity Center** — not direct IAM credentials
- ✓ Integrate your detection solution with in-scope AWS services using the permissions granted through this federation
- ✓ Configure detection, logging, and data collection as required by your solution architecture

Scope limitation: Participants will *not* have access to AWS Organization-level permissions management. You cannot view or modify Organization-level policies or SCPs.

ON-PREMISES-SIMULATED COMPONENTS

Separate AWS Organization — Standard Access Methods

Components simulating on-premises infrastructure are hosted in a separate AWS organization. Participants access these hosts using standard enterprise remote access methods.

- ✓ Standard access methods provided: **OpenVPN, SSH, and RDP** to on-premises-simulated hosts
- ✓ AWS **Traffic Mirroring** support is available and can be enabled for network-level telemetry collection — participants may use this capability or bring their own network monitoring approach
- ✓ Additional network components may be attached to the range's **transit gateway** to support participant infrastructure requirements

If using Traffic Mirroring, configure your integration before your execution week. See AWS Traffic Mirroring documentation for setup details. Contact the evaluation team if you need transit gateway attachment support.

🏗️ **Hybrid architecture:** The Scenario 2 Detection range and the Protections range are each hybrid enterprise deployments spanning **two independent AWS organizations** — one representing cloud-native infrastructure, one simulating on-premises infrastructure. Your solution must integrate with both: cloud-native components via IdP-federated AWS account access, and on-premises-simulated components via standard access methods (OpenVPN, SSH, RDP). Solutions that integrate with only one environment will have coverage gaps against techniques executed in the other.

REQUIREMENTS

Participation Requirements

Eligibility

REQUIREMENT	DETAILS
General availability	Product must be GA with paying customers at time of evaluation execution
Technical support capacity	Vendor must be able to provide dedicated technical support during the evaluation period (Aug–Oct 2026)
Results publication consent	Agreement to public results publication on the ATT&CK Evaluations website in December 2026
Citation policy compliance	Commitment to follow MITRE citation policy when referencing results in marketing and sales.

What You Must Provide

1. ### Vendor Tooling Declaration

   Declare applicable product category tags: EDR, XDR, Cloud Security, Identity Protection, MDR Service, Network Detection, Managed Service / MSSP, SIEM Integration, Endpoint Protection. Multiple tags permitted.

2. ### Configuration Documentation

   Document your solution's deployment architecture, configuration settings, and integration requirements within the test environment. Submitted prior to evaluation execution.

3. ### Technical Resource Assignment

   Assign dedicated technical contact(s) who can provide support, answer questions, and assist with integration during the evaluation execution period. Contacts must be reachable within 4 hours during execution weeks.

4. ### Draft Results Review

   Review draft results for technical accuracy before public publication. This is your opportunity to verify configuration and testing correctness. Review window and timeline TBC.

5. ### Citation Policy Compliance

   Follow MITRE citation policy guidelines when referencing results in marketing, sales conversations, and public communications.

model.

📋 **Draft results review window: November 9–20, 2026.** During this period you may review your draft results for technical accuracy, submit backend detection timestamp data, and flag discrepancies before public publication in December 2026. The dispute process and submission instructions will be communicated during onboarding.

PREPARATION GUIDANCE

How to Prepare Before Your Evaluation

These preparation steps will materially improve your TES scores. Start before your execution week — some require architecture decisions that take weeks to implement.

Detection Coverage (DC) DQI

DC is scored at the behavior level — one DC score per ATT&CK technique regardless of how many alerts you generate.

- ✓ Ensure alerts include all 7 context elements: WHO, WHAT (ATT&CK mapped), WHEN, WHERE, HOW, WHY, Recommended Action
- ✓ DC-3 requires explicit ATT&CK technique mapping — tactic-level only scores DC-2
- ✓ Test your alert quality against specific ATT&CK techniques in Scenario 1 before your execution week

Detection Precision (DP) DQI

DP rewards consolidation. One incident case covering 50 behaviors scores far higher than 50 separate cases for the same behaviors.

- ✓ Configure your correlation rules to consolidate related adversary activity into unified incident cases
- ✓ Test false positive rate against the business activity types used in FP validation testing
- ✓ Reduce case fragmentation — the case-count penalty is significant

Detection Speed (DS) DQI

DS measures time from technique execution to first automated platform alert. Sub-15-minute = Real-Time tier (score: 1.0).

- ~~Identify and reduce pipeline bottlenecks (API polling intervals, batch ingestion delays)~~

Protection Coverage (PC) & Timing

PQI

Block timing is the primary PC tier determinant. Blocking before the first critical stage scores PIW 1.0× (maximum); blocking at the terminal stage scores PIW 0.25× (minimum).

- ✓ Block alerts should also include the 7 context elements — full context at Stage 1–2 → PC-3; partial → PC-2
- ✓ If using "observe before block" strategy, document it in advance with supporting telemetry plan
- ✓ Test block triggers against the specific technique categories in Scenario 1

📋 The Full Methodology Specification contains complete scoring formulas, all 7 context element definitions, tier assignment matrices, and the cloud latency timestamp submission process. Reading it before your execution week is the highest-leverage preparation activity.

ONBOARDING PROCESS

What Happens After You Sign

1 **Contract Processing & Confirmation**

MITRE confirms receipt of signed contract and provides onboarding documentation. Group 1 lottery participants receive lottery results after March 6 deadline.

2 **Execution Week Selection**

Group 1 lottery winners select from available execution weeks (preferred order). Standard registration vendors select on first-come, first-served basis after March 7.

3 **Technical Kickoff — June 2026**

MITRE evaluation team conducts technical kickoff with your assigned technical contacts. Covers environment access, integration requirements, and communication protocols for the evaluation period. Configuration documentation is due

5 months before your execution week

.

- **Week 1–2 — Tool Installation:**
  Vendor deploys and configures their solution into the test environment.
- **Week 3 — Validation & Training:**
  Integration validation, configuration sign-off, and MITRE-led training on the vendor's product to ensure the evaluation team can accurately assess its capabilities.
- **Week 4 — Execution:**
  Adversary emulation scenarios are executed. Both detection and protection phases run during this week. Your technical contacts are available throughout.

5 **Draft Results Review — November 9–20, 2026**

MITRE provides draft results for your review. Submit backend detection timestamp data, flag technical discrepancies, and verify configuration accuracy before public publication. MITRE does not change scores based on competitive preference — only technical accuracy corrections are accepted.

6 **2026 Results & Insights — December 2026**

MITRE announces the completion of the 2026 ATT&CK Evaluations and publishes a public aggregation of insights from across the evaluation cycle. The publication highlights collective findings, defensive trends, and lessons learned while preserving the confidentiality of individual participant results.

KEY DATES

Timeline

Feb 17, 2026

Call for Participation opens.

Registration and contract submission begins.

Closed

Mar 6, 2026

Group 1 Lottery deadline (11:59pm ET).

Closed

Mar 3, 2026

Scenario 2 technical specifications published.

Full infrastructure details, integration requirements, and architecture documentation released.

Closed

Mar 7 – Apr 30

Standard registration window.

Signed contracts accepted. Execution weeks assigned on first-come, first-served basis.

Closed

Apr 30, 2026

Application deadline.

No contracts accepted after this date.

Closed

Aug – Oct 2026

Evaluation execution phase.

Adversary emulation scenarios run across all vendor execution weeks.

Now

Nov 9-20 2026

Draft results review window: November 9–20, 2026.

Participants review their private draft results, submit backend timestamp data, flag technical discrepancies, and verify configuration accuracy before results are finalized.

Upcoming

Public results & insights publication.

MITRE announces completion of the 2026 ATT&CK Evaluations and publishes aggregated insights, trends, and lessons learned from the private evaluations. Individual participants results may remain private.

Upcoming

Timeline subject to adjustment based on participation volume.

🎯 **Group 1 Lottery advantage:** Submit a complete, signed contract between February 17 – March 6, 2026 (11:59pm ET) to enter the lottery for preferred execution week selection. Preferred weeks allow more preparation time and may have fewer conflicts with other major security events. Standard registration (March 7+) assigns weeks on a first-come, first-served basis from remaining available slots.